

homework-4

```
library(bis557)
library(cas1)
library(reticulate)
#> Warning: package 'reticulate' was built under R version 3.6.2
library(stats)
use_condaenv("r-reticulate")
Sys.which("python")
#>          python
#> "/usr/bin/python"
```

1. In Python, implement a numerically-stable ridge regression that takes into account colinear (or nearly colinear) regression variables. Show that it works by comparing it to the output of your R implementation.

We build a python function for ridge regression and fit data and compare it with the previous homework's result.

```
import numpy as np

def ridge(X, y, lambda_param):
    u, s, vh = np.linalg.svd(X, full_matrices=False)
    D = np.diag(s/(s**2 + lambda_param))
    beta = np.dot(np.dot(np.dot(vh.T, D), u.T), y)
    return beta
```

This chunk generates matrix to fit in python so that the input data will be the same as R function.

```
data(iris)
df_no_na <- model.frame(Sepal.Length ~ ., iris)
py$X <- model.matrix(Sepal.Length ~ ., df_no_na)
yname <- as.character(Sepal.Length ~ .)[2]
py$y <- matrix(df_no_na[, yname], ncol = 1)
```

This chunk runs ridge regression given X, y and lambda, the penalty.

```
py_fit = ridge(X, y, .01)
```

This chunk runs previously implemented ridge regression function and compare the coefficients with python results.

```
fit_my_ride <- ridge(Sepal.Length ~ ., iris, lambda = .01)
cbind(py$py_fit, fit_my_ride$coefficients)
#>          [,1]      [,2]
#> [1,]  2.1497591  2.1497591
#> [2,]  0.5032983  0.5032983
#> [3,]  0.8271273  0.8271273
#> [4,] -0.3215299 -0.3215299
#> [5,] -0.7049677 -0.7049677
#> [6,] -0.9991592 -0.9991592
```

We could see that the function in python, using np.linalg.svd, returns to same result as the previous implemented function in R.

In the colinear case:

```

data(iris)
iris$Sepal.Width_coll <- iris$Sepal.Width*1.5+0.1
df_no_na <- model.frame(Sepal.Length ~ ., iris)
py$X <- model.matrix(Sepal.Length ~ ., iris)
yname <- as.character(Sepal.Length ~ .)[2]
py$y <- matrix(df_no_na[,yname], ncol = 1)

py_ridge_coll =ridge(X,y,.01)

py_fit <- py$py_ridge_coll
fit_my_ridge <- ridge(Sepal.Length ~ ., iris, lambda = .01)
cbind(py$py_ridge_coll, fit_my_ridge$coefficients)
#>           [,1]           [,2]
#> [1,]  2.11918470  2.11918470
#> [2,]  0.05713734  0.05713734
#> [3,]  0.82706790  0.82706790
#> [4,] -0.32169400 -0.32169400
#> [5,] -0.70444942 -0.70444942
#> [6,] -0.99850564 -0.99850564
#> [7,]  0.29762448  0.29762448

```

We could see that in the collinear case, the python function still works and has the same result as R.

2. Create an “out-of-core” implementation of the linear model that reads in contiguous rows of a data frame from a file, updates the model. You may read the data from R and send it to your Python functions for fitting.

Out-of-core refers to processing data that is too large to fit into main memory and the algorithm would allow to access data in a sequence. [Ref: <https://machinelearning.wtf/terms/out-of-core/>]

We write a function for fitting linear model in python using qr decomposition. We read the data in batches from r and fit the current batch with the python function and update the model until we fit all the data.

This chunk is the python function to fit the linear model.

```

import numpy as np

def py_lm(X,y):
    qrs = np.linalg.qr(X)
    q = qrs[0]
    r = qrs[1]
    beta = np.dot(np.dot(np.linalg.inv(r), q.T),y)

    return beta

#simulate data
n <- 1e5
X <- matrix(c(rnorm(n, 10, 1), rnorm(n, 5, 1), rnorm(n, 2, 2)),nrow = n)
y <- rnorm(n,50,1)

# create batches
batch <- 100
b_size <- n/batch
beta <- matrix(rep(0,ncol(X)*batch), nrow = batch)

for (i in 1:batch){
    #create batches

```

```

y_b <- y[(b_size*(i-1)+1):(b_size*i)]
X_b <- X[(b_size*(i-1)+1):(b_size*i),]
#read in python and fit model
py$X_b <- X_b
py$y_b <- y_b
beta[i,] <- py$py_lm(X_b,y_b)
}
#compute mean
beta_final <- apply(beta,2,mean)
beta_final
#> [1] 3.9262573 1.9856991 0.2007395

lm(y~X-1)
#>
#> Call:
#> lm(formula = y ~ X - 1)
#>
#> Coefficients:
#>      X1      X2      X3
#> 3.9266  1.9847  0.2009

```

Comparing with fitting `lm` directly, we get a pretty close result. If we don't have enough memory to hold all data for fitting, using out-of-core algorithm would be a decent choice.

3. Implement your own LASSO regression function in Python. Show that the results are the same as the function implemented in the `casl` package.

We build a LASSO function in python here. It differs the Ridge regression by the first order norm for the penalization term. Referring to [<https://stats.stackexchange.com/questions/17781/derivation-of-closed-form-lasso-solution>], we find the ridge parameters:

```

def py_lasso(X,y, lambda_param):
    n = len(X)

    qrs = np.linalg.qr(X)
    q = qrs[0]
    r = qrs[1]
    #least square
    b_ls = np.dot(np.dot(np.linalg.inv(r), q.T),y)
    #soft hold
    b_max = np.maximum(np.abs(b_ls)-lambda_param,0)
    beta = np.sign(b_ls)*b_max

    return beta

#simulate data as page 192
n <- 1000
p <- 5
X <- matrix(rnorm(n * p), ncol = p)
beta <- c(3, 2, 1, rep(0, p - 3))
y <- X %*% beta + rnorm(n = n, sd = 0.1)
bhat <- casl_lenet(X, y, lambda = 0.01)
bhat
#>           [,1]
#> [1,] 2.9945196
#> [2,] 1.9862863

```

```

#> [3,] 0.9896681
#> [4,] 0.0000000
#> [5,] 0.0000000

py$X <- X
py$y <- y
py$py_lasso(X,y,.01)
#>           [,1]
#> [1,] 2.9949377
#> [2,] 1.9873291
#> [3,] 0.9888466
#> [4,] 0.0000000
#> [5,] 0.0000000

```

```

bhat <- casl_lenet(X, y, lambda = 0.1)
bhat
#>           [,1]
#> [1,] 2.9017117
#> [2,] 1.8849432
#> [3,] 0.9095382
#> [4,] 0.0000000
#> [5,] 0.0000000

py$py_lasso(X,y,.1)
#>           [,1]
#> [1,] 2.9049377
#> [2,] 1.8973291
#> [3,] 0.8988466
#> [4,] 0.0000000
#> [5,] 0.0000000

```

Comparing lasso results with penalty .01 and .1 with casl package, we got very similar results.

4. Propose a final project for the class.

I am thinking of using CNN to classify cat and dog images from a kaggle dataset cat and dog [<https://www.kaggle.com/tongpython/cat-and-dog>]. I will follow the guidance from class and build a convolutional neural network with keras. It will take steps to train and test the model with two datasets of images, and classify the images.

I will try 2D layers with different activation functions, kernel sizes, and different batch sizes and epochs to get a model with relatively high prediction accuracy. The training process could provide visualization along with training process.